

Web Search Engines — Reviews-Rehashed

Himaja Rachakonda, Anirudhan J. Rajagopalan
N14633788, N18824115
hr970, ajr619

March 21, 2016

1 Objective

As consumers search online for product information and to evaluate product alternatives, they often have access to dozens or hundreds of product reviews from other consumers. These customer reviews are provided in addition to product descriptions, reviews from experts, and personalized advice generated by automated recommendation systems. Each of these options has the potential to add value for a prospective customer. Online reviews have become one of the most powerful factors which impact the consumer behaviour in the E-commerce industry. Both the consumers as well as the businesses gain a ton of information from the virtual voice of reviews. Online reviews transformed the phenomenon of the simple word-of-mouth feedback into a viral form of virtual feedback, which is not only bringing immense customer satisfaction in relying on them but also showing mind boggling returns in the businesses as well. Therefore, the aim of "Reviews-Rehashed" is to make the relevant reviews more accessible to the users.

Reviews-Rehashed is a specialized search engine to search product reviews pertaining to a product and its features. The search queries are in the form of $\langle \text{product}, \text{feature} \rangle$ tuples. This search engine will crawl reviews from various e-commerce websites such as Amazon and BestBuy.com to fetch reviews data of various products. It searches for the user query and presents the links to the results after ranking the documents based on the intrinsic quality of the document as well as the retrieval score. The quality of the document is based on a static score derived from the features of a review like the ratings, review-date, comments etc. We have implemented the search engine for electronics - mobiles data from Amazon.com. There is a huge scope of expanding this to other products and other domains as well.

2 Data Sources

Amazon.com and BestBuy.com were the options considered to collect information regarding products and their reviews. To build a crawler to fetch the data, a preliminary study was conducted on the following to understand the website page layout and structure —

1. User behaviour to search for a particular product and review.
2. Patterns in the URLs of the websites to reach to a particular product / product category / review / etc.
3. Examining the sitemap.xml of Amazon.com and BestBuy.com to get all the URLs of the domain.

Sitemap is a list of pages of a web site accessible to crawlers or users. It can be either a document in any form used as a planning tool for Web design, or a Web page that lists the pages on a Web site, typically organized in hierarchical fashion. Sitemaps make relationships between pages and other content

components. It shows shape of information space in overview. Sitemaps can demonstrate organization, navigation, and labeling system.

Sitemaps are a useful tool for making sites built in Flash and other non-html languages searchable. If a website's navigation is built with Flash, an automated search program would probably only find the initial homepage; subsequent pages are unlikely to be found without an XML sitemap.

This helped us to make sure all the review pages are searchable by our search-engine and also to find useful patterns which help in identifying each product and review within each of the domains of Amazon and BestBuy. Because of the vast expense of the data and the crawling required to be done, we have restricted our project to Amazon Review Data pertaining to Mobile Phones under Electronics Sections. This can be scaled easily to other categories and domains seamlessly with appropriate crawler jobs.

All the review pages of a particular product with an ASIN Id can be found at a URL in the following format —

`www.amazon.com/<Anything>/product-reviews/<ASIN>/`

Every Review of a product has a unique Review ID, and therefore each review can be accessed by a URL in the following format —

`http://www.amazon.com/gp/customer-reviews/<Review-ID>/`

3 Data Collection

Data Collection was executed in two phases —

Amazon Product Crawler: Fetches all the unique Ids for all mobile phone products from Amazon.com. This is required to construct the seedUrls for each product as mentioned above in the Section 2, Data Sources Section.

Reviews Crawler: This crawler takes the output of the the Amazon Product Crawler as the input to construct a seed Url for each product Id. All the unique IDs (ASINs) are stored, which are used to fetch Reviews from each product. These reviews can be accessed from this page through the links in the pagination bar. Each page consists of 10 reviews and the crawler goes to the depth until which there are no more review pages to be crawled for that product. Therefore, this product level URL is given as a seedURL for every product to crawl across various review pages through the links in the pagination bar.

The data has been collected for around 500 products with a total of 70000 reviews. The following data about each review is collected as indexable fields based on which the review will be scored during the search time —

1. Review Content
2. Review Title

3. Review Date
4. Number of Comments to the review
5. Number of people found this review helpful
6. Number of images posted by the user in this review
7. Ratings of the review
8. Is the product a Verified Purchase
9. Review Length
10. Review Id
11. ASIN (Product Id)

4 Search Engine Architecture

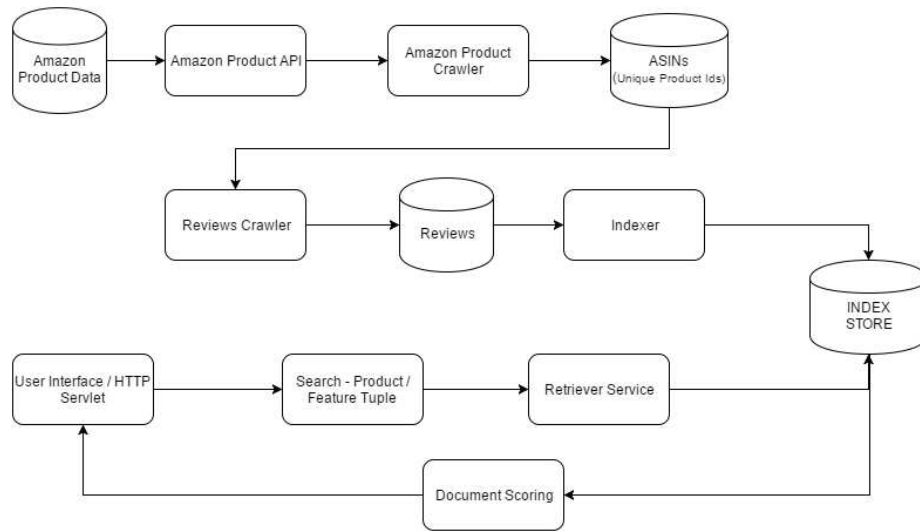


Figure 1: Reviews-Rehashed Search Engine Architecture

The search engine will be built by using the following different components.

Preprocessing module: This module takes care of preprocessing the data and dumps the processed data into Lucene.

Amazon WebCrawler: This module fetches all the ASINs (unique Product Ids) for Mobile phones under Electronics Sections.

Reviews WebCrawler: This module fetches all the reviews of the products from the product Ids retrieved from the Amazon WebCrawler module.

Document Scoring: This module implements the ranking algorithm by which the documents are retrieved in the search algorithm.

Retriever Service: This module receives the search query from the Web interface and interacts with the index store to retrieve results. The retrieved results, based on ranking algorithm in Doc Scoring module, are returned back to the web interface.

Web interface: A web interface for the user to issue the search queries. This will be an interface with two text boxes. One for the product name and another for the feature name.

Index store: This is a Lucene index store, an inverted index, which maintains all indexable data required while retrieving search results. The data will be populated and updated periodically by a crawler job that looks for updates to the data. The data will be stored in lucene after preprocessing.

5 Document scoring

The documents are scored based on a personalized page rank algorithm which takes the intrinsic quality of the document into consideration as well as the lucene score generated. In many search engines, we have available a measure of quality $g(d)$ for each document d that is query-independent and thus static. This quality measure may be viewed as a number between zero and one. The net score for a document d is some combination of $g(d)$ together with the query-dependent score induced by lucene. This kind of ranking algorithm demands an accumulation of an evidence of a document's relevance from multiple sources.

The personalized pageRank algorithm gives 25% weightage to the quality of the document and 75% weightage to the lucene score of the document. The quality of the document is again a combination of weights given to each of the dependent factors.

The pageRank algorithm is as follows —

$$pageRank(d) = 0.25 * g(d) + 0.75 * (lucenescore) \quad (1)$$

The quality $g(d)$ of the documents is as follows

$$g(d) = 0.6 * H + 0.20 * V + 0.13 * C + 0.07 * I,$$

where H : = Number of people who found the review helpful,

V : = The purchase is a verified purchase,

C : = Number of comments to the review,

I : = Number of images in the review

As all the dependent factors and the lucene score are unbounded non-negative variables with different frequency distributions, it is necessary to normalize them before we use them in the pageRank algorithm in equation (1) to have consistent scores. A sigmoid mathematical function is chosen to normalize these unbounded variables with a multiplicative factor of the x-axis to obtain scaled values between 0 and 1, based on distributions of each of these variables.

A sigmoid function is a bounded differentiable real function that is defined for all real input values and has a positive derivative at each points range is always between 0 and 1.

$$S(t) = \frac{1}{1 + e^{-t}}. \quad (2)$$

The following graph shows the spread of the sigmoid function for various multiplicative factors —




Figure 2: sigmoid variations with different multiplicative factors

6 Experiments and Observations

The reviews of the products in the search results are compared to the top reviews of those products in Amazon.com. We observed that 80% of the times all the search results match with the top reviews of the corresponding products in Amazon. The cases where it performs well and where it doesnot perform well are discussed below—

Positive cases: 1 . The best results were obtained with the final scoring formula discussed in the Section 5. This works pretty well for most of the queries with valid product name and feature name.

Negative cases: 1 . The results are not very accurate when conjunctive queries are given in the feature. An example can be "Camera Battery" as a

Review Rehashed Search product reviews

Feature * Product Group

Fraud
 I have been deceived since I bought a NEW Iphone and they sent me a USED, BROKEN & with missing parts device. Since I live in Argentina I had the phone sent to a friends house in Miami and when she came to Argentina she brought me the phone. I was surprised when I found out that the cable USB/charger was broken. However, I used the one from my IPAD that perfectly fits. After a week, the battery stopped charging, so I took it to a Mac store in my country. I was told that the phone was not new device (it was refurbished) and further, had missing parts inside. Unfortunately the 30 days return time has elapsed. Therefore, it is impossible for me to contact the seller and claim for the restitution of my money.
<http://www.amazon.com/gp/aw/review/B01GNNXSGR2/2WBPB6OE03I>

This was a refurbished product and iam satisfied with it
 This was a refurbished product and iam satisfied with it. I ordered it for my daughter and she is happy. There is not a lot of storage but it's great for her. The battery life is good, there was no visible damage to the outside or the screen. It came with a charger. The packaging was not very good at all. It arrived in an oversized box flopping around loosely with only a thin piece of plastic to protect the screen. For the money I believe it's a good deal. We've not had any problems with it so far. We also ordered a otter box defender for it. It was cheap and has done its job already. Well worth the money.
<http://www.amazon.com/gp/aw/review/B00YV5GQU4/R2/CMGUMTFZYU7>

This was a refurbished product and iam satisfied with it
 This was a refurbished product and iam satisfied with it. I ordered it for my daughter and she is happy. There is not a lot of storage but it's great for her. The battery life is good, there was no visible damage to the outside or the screen. It came with a charger. The packaging was not very good at all. It arrived in an oversized box flopping around loosely with only a thin piece of plastic to protect the screen. For the money I believe it's a good deal. We've not had any problems with it so far. We also ordered a otter box defender for it. It was cheap and has done its job already. Well worth the money.
<http://www.amazon.com/gp/aw/review/B00YV5E1UR/2/CMGUMTFZYU7>

Figure 3: Scoring with different multiplicative factors for the dependent factors of the score - numHelpful=0.05, isVerifiedPurchase=100, numComments=0.1, numImages=1

search query for the feature of a product. We expect to see reviews containing both camera and battery in the review, however we see some reviews containing only battery and some reviews down the list, containing camera. This behaviour is erratic and the search results are not always consistent

2. Search results are better without the scoring algorithm when the product query is null. For example, if we just give "Great phone, good battery life" for feature, search results did not show results which have the exact same line in the review title. This is because of the other factors we consider for determining the static score/quality of the document. However, in this scenario, we also expect to see some direct matches with reviews as well.

7 Challenges

Understanding patterns from sitemap.xml was a time — taking task.

Manually checking the weights was a tedious task, evaluation was difficult.
 Have to train the data to tune these weights. (future work)

8 Future work

NLP module : This will essentially be a machine learning model that will be pre trained using the data collected by the crawler. We will use this to

rank and retrieve the list of all matching sentences for a given query. This module can be developed to have the intelligence to do the following tasks–

Summary Extraction: A summary from the top - K reviews is presented to the user, giving an overall picture of the opinion of the crowd on this product’s feature.

Opinion Mining : A module which shows the positive and negative reviews of a product

Data Analytics : A module to collect data for analysis of the consumer behaviour in learning about a product and / or its features.

Tuning Score parameters: Score Parameters can be tuned to more accurate values by training the data and running models, to improve the current scoring algorithm. Machine Learning to Rank the documents based on these parameters may give a better performance with respect to the relevancy of the search.

9 To DO

references

proof reading

hyper references

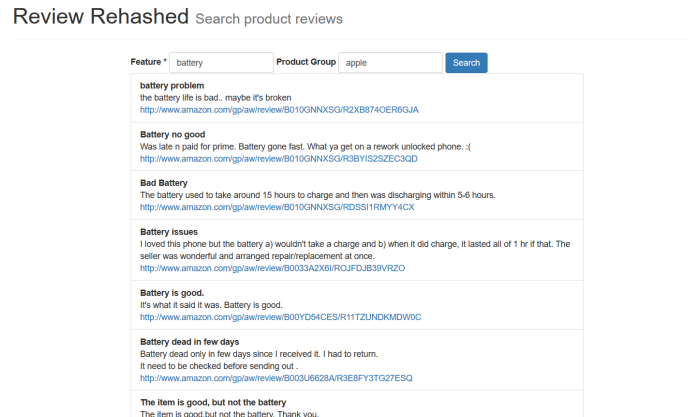


Figure 4: Default lucene scoring with feature query boost of 0.75 for title and 0.25 for content.

Review Rehashed Search product reviews

Feature *	battery	Product Group	apple	Search
battery problem the battery life is bad... maybe it's broken http://www.amazon.com/gp/aw/review/B010GNNXSG/R2XB874OER8GJA				
Battery is good. It's what it said it was. Battery is good. http://www.amazon.com/gp/aw/review/B00YD54CES/R11TZUNDKMDW0C				
Battery no good Was late n paid for prime. Battery gone fast. What ya get on a rework unlocked phone. :(http://www.amazon.com/gp/aw/review/B010GNNXSG/R3BYIS2S2EC3QD				
Bad Battery The battery used to take around 15 hours to charge and then was discharging within 5-6 hours. http://www.amazon.com/gp/aw/review/B010GNNXSG/RDSS11RMY4CX				
Battery issues I loved this phone but the battery a) wouldn't take a charge and b) when it did charge, it lasted all of 1 hr if that. The seller was wonderful and arranged repair/replacement at once. http://www.amazon.com/gp/aw/review/B003A2X8I/R0JFDJB39VRZO				
The item is good, but not the battery The item is good, but not the battery. Thank you. http://www.amazon.com/gp/aw/review/B003A2X8I/R29TL0ENLW6OP				
No scratches Battery life is No scratches Battery life is good				

Figure 5: Default lucene scoring without boosting.

Review Rehashed Search product reviews

Feature *	battery	Product Group	apple	Search
Battery issues I loved this phone but the battery a) wouldn't take a charge and b) when it did charge, it lasted all of 1 hr if that. The seller was wonderful and arranged repair/replacement at once. http://www.amazon.com/gp/aw/review/B003A2X8I/R0JFDJB39VRZO				
battery problem the battery life is bad... maybe it's broken http://www.amazon.com/gp/aw/review/B010GNNXSG/R2XB874OER8GJA				
Bad Battery The battery used to take around 15 hours to charge and then was discharging within 5-6 hours. http://www.amazon.com/gp/aw/review/B010GNNXSG/RDSS11RMY4CX				
Battery is good. It's what it said it was. Battery is good. http://www.amazon.com/gp/aw/review/B00YD54CES/R11TZUNDKMDW0C				
Battery no good Was late n paid for prime. Battery gone fast. What ya get on a rework unlocked phone. :(http://www.amazon.com/gp/aw/review/B010GNNXSG/R3BYIS2S2EC3QD				
Battery dead in few days Battery dead only in few days since I received it. I had to return. It need to be checked before sending out. http://www.amazon.com/gp/aw/review/B003J6628A/R3E8FY3TG27ESQ				
The item is good, but not the battery The item is good but not the battery. Thank you.				

Figure 6: When the normalization weights are equal.

Review Rehashed Search product reviews

Feature * Product Group

warning do not update to iOS 9.2

I did love this iPhone however, I had to replace the battery twice, and ultimately I can no longer use this iPhone because I let it update to iOS 9.2 and it froze my phone to recovery mode with the iTunes logo. I chatted with 3 techs on apple.com to try to resolve the issue, and from 11pm to 4am, this iPhone would not recover back to normal, so I took it to an apple store later in the day, and to my sadness, the phone could not be restored, they offered to sell me another iPhone being that this phone was pass warranty, the replacement phone would have cost way more than the price that I paid for this phone. I am very upset, because I just recieved a new battery for the phone last week, and I had just received an iHome that I purchased the same day that this happened, and here is the kicker... I had just put minutes on this phone 5 minutes before letting it update.... So dissapointed

<http://www.amazon.com/gp/aw/review/B007VOF51M/R2MG6D3YY3NWWY1>

I Would Not Recommend Buying an iPhone

I owned an iPhone 4 for more than a year. Prior to getting the phone, I had owned an old standard phone without Internet. When one of my co-workers first showed me the iPhone and what it could do, I could not believe it and I went out and purchased one the next day without even investigating other phones. When I first got it, I loved it. It was amazing, almost like having a computer in your pocket. I was in love with the phone. Then, over time, as I started using the phone more to go to various websites, I realized that many websites that I was going to had blank spots where there were supposed to be pictures, or the sites did not work at all. I just assumed that it was not possible to view these sites on a smart phone. I thought it was a limitation of the technology. Then when some of my students would show me their smart phones, I noticed that they were able to view the same websites that I was not. I wondered why that was, and what was in common between all those phones that could view those websites was that they were not iPhones. I came to realize that Apple does not allow flash player on their iPhones, or their iPads for that matter. This upset me quite a bit. I read Steve Jobs' explanation in a long letter that he posted online. It seems to me that he simply had a

Figure 7: This scoring is weird. But looks decent.